

Advanced Operator / Codex Track

Advanced training for power users who want to supervise Codex and similar local full-control systems with judgment, testing, and rollback habits.

Format	Private coaching or cohort sessions with local practice, repo walkthroughs, and verification labs.
Level	Level 4: advanced operator and local-agent supervision
Audience	Technically curious professionals, builders, content-system owners, and advanced AI users ready for local controlled environments.

Learner outcomes

- Understand what local full-control systems can access and why that changes the supervision model.
- Use repo inspection, file edits, tests, browser checks, Git commits, and deployment review responsibly.
- Develop rollback thinking so AI-assisted work remains reversible, inspectable, and owned by the human.
- Operate Codex-style systems with explicit boundaries, verification habits, and human approval points.

Guiding questions

- What changes when an AI system can read files, run commands, inspect browsers, and edit projects?
- Which actions are safe for an agent to take, and which require explicit human approval?
- How will we verify outputs with tests, screenshots, diffs, and deployment checks?
- How will we recover if an AI-assisted change is wrong?

Readiness checks

- The operator can explain what local files, commands, browsers, and deployments the agent can affect.
- The operator can request scoped changes and inspect diffs before accepting them.
- The operator can run appropriate verification before committing or publishing.
- The operator has a rollback path and knows when to stop for human review.

Research-backed adoption practices

- Use least-privilege thinking: give local agents only the access, scope, and autonomy the task actually requires.
- Assume agent outputs and actions need verification through diffs, tests, screenshots, logs, or independent review.
- Treat prompt injection, sensitive information disclosure, and excessive agency as design risks for tool-using systems.
- Keep rollback, audit trail, and human approval points visible before using full-control agents on higher-risk work.

Curriculum modules

From chatbot to local operator

We compare ordinary LLM sessions with local agents that can read files, run commands, inspect browsers, and edit projects.

Supervising code and content changes

Participants learn to ask for repo inspection, scoped edits, diffs, tests, and clear summaries before trusting changes.

Full-control safety boundaries

We define where local access helps, where it becomes risky, and what should require explicit human approval.

Deployment and rollback practice

Participants practice commit hygiene, public-site checks, deployment verification, and recovery paths.

Operator judgment

Advanced users learn when to let the agent proceed, when to require confirmation, when to inspect manually, and when to stop.

First prompting lab

These short starter activities help learners experience prompt design before longer workflow projects. They are adapted from the activity patterns researched on Everybody Prompts and similar beginner practice formats.

Ask-for-a-prompt warmup (Pre-prompting)

Learners discover that prompting can begin by asking the model to help shape the prompt itself.

Build: A reusable prompt card with task, audience, context, constraints, output format, and review criteria.

Safe input: Use a public topic, fictional scenario, or harmless personal learning goal.

Starter prompt: I want to use AI for [safe topic or task], but I am not sure how to prompt well. Ask me up to five clarifying questions, then give me three prompt options: quick, structured, and advanced. Include what each prompt is good for and what I should check afterward.

Structured research table (Basic research prompting)

Students learn to ask for organized output instead of accepting a long, shapeless answer.

Build: A comparison table with columns the learner specified, plus a list of claims to verify.

Safe input: Use a public topic such as renewable energy, storage devices, local history, or a general product category.

Starter prompt: Create a beginner-friendly research table about [public topic]. Use columns for concept, short explanation, practical use, common confusion, and what to verify in reliable sources. After the table, suggest three better follow-up questions.

Timeline builder (Chronological organization)

Students see how prompt format changes understanding by turning a topic into sequence, development, and context.

Build: A chronological timeline with significance notes and uncertainty flags.

Safe input: Use public historical, technical, cultural, or organizational topics.

Starter prompt: Create a chronological timeline for [public topic]. Include year or period, event, why it mattered, what changed afterward, and confidence level. Mark anything that needs source checking.

Email tone trio (Business email drafting)

Learners experience AI as a writing partner while keeping voice, promises, and final judgment human-led.

Build: Three email versions, a tone comparison, and a send-before-review checklist.

Safe input: Use a fictional or sanitized email scenario with no customer, patient, employee, financial, or proprietary details.

Starter prompt: Using this fictional email scenario, draft three versions: concise, warm, and more formal. Explain what changed in each version, recommend one, and list facts, promises, names, dates, and tone choices a human should confirm before sending: [scenario].

Image observation drill (Image description)

Students distinguish visible evidence from inference, uncertainty, and possible overclaiming.

Build: A two-column observation table separating what is visible from what is inferred.

Safe input: Use a public-domain, stock, classroom, or non-sensitive image. Avoid faces or private settings unless permissions are clear.

Starter prompt: Describe this image for a careful learner. Separate visible details from inferences. Add a table with observation, evidence in the image, confidence, and what should not be assumed. Finish with three questions a human reviewer should ask.

Career fit map (Career research)

Learners practice giving context, asking for options, and requesting useful decision criteria.

Build: A career option table with fit reasons, education paths, tradeoffs, and next questions.

Safe input: Use a fictional profile or a sanitized personal summary that excludes private employer data, compensation, references, and confidential projects.

Starter prompt: Using this safe career profile, suggest career directions in a table with role, why it may fit, skills to build, typical education or training path, tradeoffs, and next research questions. Do not make the decision for me: [profile].

Proof-bank resume drill (Resume creation)

Students learn to turn scattered experience into defensible claims and then ask for critique.

Build: Six resume bullets, proof bank, critique table, and claims-to-evidence checklist.

Safe input: Use a fictional or sanitized work history without private employer details, references, compensation, or confidential outcomes.

Starter prompt: Turn this sanitized work history into a proof bank and six resume bullet options. For each bullet, list the evidence needed to support it, possible overstatement risk, and one stronger revision. Flag anything that sounds inflated: [work history].

Brand connotation check (Branding feedback)

Learners see how AI can brainstorm associations while human judgment still checks audience, culture, and evidence.

Build: A naming or messaging table with positive associations, risks, audience fit, and questions to test.

Safe input: Use fictional brand names, public examples, or early ideas that are not confidential.

Starter prompt: Evaluate these fictional brand or project names for connotations. Create a table with name, positive associations, possible drawbacks, audience fit, words to avoid, and questions we should test with real people: [names].

Survey cleanup and format (Survey construction)

Students learn to ask for editing, consistency, and application-ready formatting in one clear request.

Build: A cleaned survey list plus a format-ready version for a form tool or spreadsheet.

Safe input: Use fictional survey items or public workshop feedback questions. Avoid private personnel, patient, customer, or sensitive demographic data.

Starter prompt: Clean up these fictional survey items. Fix grammar, remove leading language, make the scale consistent, and return two outputs: a numbered review table and a form-ready list in the format [Required] [Question type] Question text: [items].

Chart storyboard (Chart creation)

Learners separate the data question, chart type, audience, and caveats before asking for a visual.

Build: A chart plan with data needed, recommended chart type, annotation ideas, and cautions.

Safe input: Use public data, fictional data, or a small hand-written sample. Do not use sensitive spreadsheets in public tools.

Starter prompt: Help me plan a chart for [public or fictional data topic]. Recommend chart type, needed columns, sample data structure, title, labels, annotations, caveats, and three checks before presenting the chart.

Cloze quiz maker (Quiz generation)

Students practice precise output instructions, distractor quality, and iterative repair.

Build: A short cloze quiz in table format with answer key, distractors, and a quality check.

Safe input: Use a public article, textbook excerpt you have rights to use, or original instructional text.

Starter prompt: Create a five-item cloze quiz from this public or original text. Return a table with sentence, missing term, three plausible distractors, correct answer, and explanation. Then critique the quiz for ambiguity and weak distractors: [text].

Lesson or page outline (Lesson plans and website pages)

Learners see how AI can structure material into a usable outline before any polished writing begins.

Build: A one-hour lesson plan or one-page website outline with sections, examples, activities, and review notes.

Safe input: Use public topic notes, fictional program descriptions, or sanitized learning goals.

Starter prompt: Turn this safe topic into a one-hour lesson plan or one-page website outline. Include audience, learning goals, section sequence, examples, activity ideas, plain-language explanations, and review questions. Mark assumptions and missing information: [topic notes].

Session flow

Orient

Compare chat-based AI with local agent systems that can act inside a project environment.

Inspect

Read project structure, current state, dependencies, and risk before asking for changes.

Constrain

Write scoped instructions that define files, boundaries, verification, and approval points.

Verify

Use linting, tests, screenshots, diffs, and manual review to decide whether work is acceptable.

Recover

Practice commit hygiene, rollback thinking, and post-deployment checks.

Practice labs

Repo inspection memo

Build: a pre-edit memo naming app structure, relevant files, dependencies, git state, risks, and proposed smallest change. Use: a low-risk local practice repo. Do: inspect before editing. Check: the operator understands the project surface.

Scoped edit diff

Build: a tiny change set plus a diff review note explaining what changed, what did not, and why. Use: a small public content or UI change. Do: enforce non-goals and reject unrelated churn. Check: scope stayed tight.

Verification evidence packet

Build: commands run, routes checked, screenshots or render checks, failures, fixes, and residual risk. Use: a completed low-risk change. Do: run checks and inspect affected routes. Check: evidence supports a commit decision.

Rollback note

Build: a note naming commit, changed files, deployment check, reversal path, and confirmation evidence. Use: a practice commit or low-risk branch. Do: reason about rollback before higher-risk work. Check: reversibility is visible.

Approval boundary matrix

Build: a matrix classifying reads, edits, tests, installs, API calls, credentials, deploys, deletions, and payments by approval level. Use: fictional agent actions. Do: classify and write boundary prompts. Check: user control remains clear.

1-2 hour utility projects

These short projects are designed to help individuals experience practical AI utility while keeping inputs safe and human review visible.

How to choose a strong first project

- Real work, safe input: choose a task learners recognize, then substitute public, fictional, or sanitized material.
- Named walk-away artifact: produce a file, table, checklist, brief, agenda, glossary, recipe card, or other object.
- Visible human judgment: require the learner to check claims, reject weak output, revise tone, name uncertainty, or decide what stays human-led.

- Transferable next use: teach a repeatable pattern such as adding context, requesting structure, comparing alternatives, verifying claims, documenting the recipe, and trying it again.

Public Article Learning Kit (60-75 minutes)

Turn one public article into a tutor session and a usable study packet.

Walk-away artifact: A one-page explainer, eight-term glossary, five-question quiz, misconception list, and verification table.

Safe input: Use a public article, help page, announcement, or topic summary. Avoid private client, patient, employee, financial, legal, or proprietary material.

- Paste or summarize one public source and state your current knowledge level.
- Ask for an explainer, glossary, examples, and common misconceptions.
- Create a five-question quiz and answer it before looking at the answer key.
- Mark every claim that needs source checking before you share or rely on it.

Starter prompt: Use this public source to create a learning kit for a smart beginner: [public source or topic]. Produce a one-page explainer, eight-term glossary, three examples, common misconceptions, a five-question quiz with answer key, and a verification table listing claims I should check in reliable sources.

Decision Brief in a Box (75-90 minutes)

Use AI to structure a decision without handing the decision away.

Walk-away artifact: A one-page decision brief with criteria, option table, assumption register, pre-mortem, and next-evidence checklist.

Safe input: Use a personal, public, fictional, or sanitized decision. Keep confidential finances, contracts, medical, legal, or customer details out of public tools.

- Name two or three options, the decision deadline, constraints, and what success would look like.
- Ask for a weighted criteria table and force the model to show its assumptions.
- Run a pre-mortem for the most tempting option and a second pass for the safest option.
- Write three human-only judgment questions you will answer before acting.

Starter prompt: Help me build a decision brief without deciding for me: [safe decision context]. Include weighted criteria, an option comparison table, assumptions, likely failure modes, missing evidence, who should be consulted, and three final human judgment questions.

Meeting-to-Momentum Kit (60-90 minutes)

Convert a rough meeting idea into a focused agenda and follow-through packet.

Walk-away artifact: A timed agenda, participant prep note, decision log template, action-item table, and follow-up email draft.

Safe input: Use fictional or sanitized meeting context. Do not paste private attendee notes, personnel details, patient data, customer records, or proprietary plans into public tools.

- Describe the meeting purpose, attendees by role, time available, and desired output.
- Ask for a timed agenda with decisions, discussion prompts, and parking-lot items.
- Create a follow-up template with owners, due dates, unresolved questions, and confirmation language.
- Delete or rewrite anything that sounds like a promise, personnel judgment, or confidential detail.

Starter prompt: Using this fictional or sanitized meeting context, create a meeting kit: timed agenda, prep questions by role, likely tensions, decision log, action-item table, and follow-up email draft. Mark assumptions and anything a human should confirm before sending: [context].

Before-and-After Writing Studio (60 minutes)

Use AI as an editor, not a ghostwriter, and leave with a cleaner real draft.

Walk-away artifact: A revised email, memo, bio, announcement, or article with a change log, tone rationale, and final review checklist.

Safe input: Use a rough draft that contains no sensitive personal, client, patient, employee, legal, financial, or proprietary details.

- Paste a safe rough draft and name the audience, purpose, tone, and length limit.
- Ask for two alternative revisions: one clearer and one warmer or more concise.
- Have AI critique the revision for unsupported claims, missing context, and awkward tone.
- Make the final edits yourself and keep a before-and-after note on what improved.

Starter prompt: Revise this safe rough draft for [audience] with a [tone] tone and a [length] limit. Give me two versions, explain the changes, then critique the stronger version for clarity, unsupported claims, missing context, and anything I should verify before sending: [safe draft].

Research Launch Board (90-120 minutes)

Plan the research before the rabbit hole opens.

Walk-away artifact: A research board with question clusters, search strings, source ladder, opposing views, red flags, and claim-check table.

Safe input: Use a public topic or general business question. Do not ask AI to invent citations or make final claims without sources.

- Start with one broad public question and ask AI to break it into subquestions.
- Generate search strings, source categories, keywords to avoid, and missing perspectives.
- Ask for a claim-check table with columns for claim, source needed, confidence, and next action.
- Use the board to search actual sources; do not treat the model output as evidence.

Starter prompt: Help me build a research launch board for [public topic]. Create question clusters, search strings, source types, likely blind spots, opposing perspectives, red flags, and a claim-check table. Do not invent citations or make final claims without sources.

Workflow Recipe Card (90-120 minutes)

Turn one repeated task into a reusable recipe your future self can follow.

Walk-away artifact: A recipe card with trigger, safe inputs, prompt sequence, review gates, owner, time saved estimate, and stop signs.

Safe input: Choose a recurring task and describe it generally or with fictional examples. Keep sensitive operational data out of public tools.

- Write the current task steps from trigger to final output.
- Mark which inputs are safe, which must be generalized, and which should not enter a public tool.
- Ask AI to propose a prompt sequence and review checkpoints for the safest useful slice.
- Test the recipe with one fictional example and revise the steps until they are repeatable.

Starter prompt: Turn this recurring task into a safe AI-assisted workflow recipe: [task]. Include trigger, user, safe inputs, prompt sequence, expected output, review checkpoints, human owner, estimated time saved, risks, stop signs, and when not to use AI.

Professional Proof Portfolio (75-90 minutes)

Use AI to turn scattered experience into concrete positioning you can defend.

Walk-away artifact: A 150-word bio, six resume bullets, proof bank, interview answer outlines, and claims-to-evidence checklist.

Safe input: Use a sanitized career summary or fictionalized role description. Do not include private employer data, references, compensation, or confidential projects.

- Write a sanitized inventory of roles, strengths, projects, and outcomes you can discuss publicly.
- Ask AI to draft positioning options and separate claims from evidence.
- Create resume bullets and interview outlines, then remove anything inflated or unverifiable.
- Practice three interview questions and revise the language until it sounds true in your own voice.

Starter prompt: Using this sanitized professional summary, help me build a proof portfolio. Draft a 150-word bio, six resume bullet options, a proof bank, three interview answer outlines, and a claims-to-evidence checklist. Flag anything that sounds inflated or needs stronger evidence: [summary].

Seven-Day Action Plan (60-75 minutes)

Turn a vague goal into a week of small, scheduled actions.

Walk-away artifact: A seven-day plan with calendar blocks, next actions, obstacle responses, accountability note, and end-of-week review.

Safe input: Use personal goals and constraints that are not private, medical, financial, legal, or deeply sensitive.

- Describe the goal, available time, constraints, and what usually derails progress.
- Ask for a one-week plan with daily actions no longer than 45 minutes.
- Have AI identify likely obstacles and create if-then responses.
- Put the first two actions on your calendar and delete anything unrealistic.

Starter prompt: Help me turn this safe goal into a seven-day action plan: [safe goal]. Ask up to five clarifying questions if needed, then create calendar blocks, daily next actions under 45 minutes, obstacle responses, an accountability note, and an end-of-week review checklist.

Tangible cases

Website content update

Situation: A site owner wants to add a new section to a public website while preserving design patterns and avoiding unrelated changes.

Learner task: Ask the agent to inspect the repo, write a pre-edit memo, implement a scoped change, verify it, and produce a diff summary.

Starter prompt: Inspect this project first and write a pre-edit memo with relevant files, existing patterns, risks, and verification plan. Then add a new public content section about [topic] using existing design patterns. Do not change unrelated files. After editing, run checks, inspect desktop and mobile layout if relevant, and summarize exact files changed.

PDF companion generation

Situation: A curriculum owner wants PDF companions regenerated from site content and visually checked.

Learner task: Use the local generator, render representative PDF pages, inspect layout, record evidence, and keep public and output copies synchronized.

Starter prompt: Update the PDF companion content for [track]. Regenerate the PDFs, render representative pages to images, check for clipping, awkward spacing, missing sections, and stale content, then report output files, evidence inspected, fixes made, and residual risks.

Low-risk bug fix

Situation: A local app has a visible layout bug on mobile after new content was added.

Learner task: Have the agent reproduce the issue, inspect CSS, make the smallest targeted fix, run checks, and attach mobile verification evidence.

Starter prompt: Reproduce the mobile layout issue at [route]. Identify the smallest CSS change that fixes it, avoid unrelated refactors, run lint and build, verify at 390px and desktop width, and report the before/after evidence and any residual risk.

Prompt library

Inspect before editing

Before making changes, inspect the repo structure, relevant files, existing patterns, current git state, dependencies, and likely risks. Produce a pre-edit memo with the smallest implementation plan, expected files, checks to run, and approval triggers for deployment, credentials, deletion, installs, payments, or external API calls.

Scoped edit

Make only the requested change: [change]. Follow existing patterns, avoid unrelated refactors, and list assumptions. After editing, provide a diff summary, files changed, checks run, screenshots or render evidence if relevant, and anything a human should inspect.

Verification plan

Create and run a verification plan for this change. Include static checks, build, affected routes, desktop and mobile browser checks, screenshot or render checks where relevant, known limitations, and residual risks a human should review.

Rollback thinking

Before publishing, write a rollback note. Identify the commit or branch, changed files, deployment checks, safest reversal path, data or content that would need special care, and the exact evidence we should confirm after release.

Materials to develop

- Local agent supervision checklist
- Git and deployment primer
- Verification lab worksheet
- Rollback and recovery checklist
- Approval-boundary matrix
- Codex prompt patterns for scoped work

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- Agent control loop: inspect, edit, test, review, commit
- Local full-control safety boundary map
- Approval boundary matrix for local agent actions

Follow-up work

- Run one scoped agent-assisted change on a low-risk project.
- Capture the prompt, diff, checks, and final decision in a practice log.
- Build a personal checklist for future Codex-style work before using it on higher-risk projects.

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.