

AI Fluency Essentials

A calm, practical introduction for people who want to understand LLMs, use them safely, and build confidence with ordinary work tasks.

Format	One 90-minute Zoom session with guided practice and a follow-up handout.
Audience	Curious professionals, mixed-comfort teams, and organizations that need a shared starting point before deeper tool decisions.

Learner outcomes

- Explain what LLMs are good at, where they fail, and why human review matters.
- Use a basic prompting loop for drafting, summarizing, planning, and asking better questions.
- Recognize sensitive-data boundaries and safer ways to practice with non-sensitive examples.

Curriculum modules

The practical mental model

Participants learn how prompts, context, instructions, examples, and review shape an AI session without needing technical background.

Prompting as conversation design

We practice turning vague requests into useful context, constraints, examples, and review questions.

Safe experimentation

The session establishes clear boundaries around patient, customer, employee, legal, financial, and proprietary information.

Confidence routines

Participants leave with small repeatable habits: draft, critique, revise, verify, and decide what still needs human judgment.

Materials to develop

- AI fluency starter guide
- Prompt pattern quick sheet

- Sensitive-data boundary checklist
- Output review checklist

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- LLM session loop: prompt, context, output, review, revision
- Safe-data boundary map for public AI tools

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.