

Monthly AI Office Hours

A recurring learning forum for questions, tool changes, workflow practice, and steady coaching as AI habits mature.

Format	Monthly 60-minute Zoom sessions with question collection and lightweight follow-up notes.
Level	Ongoing level: sustained fluency and habit support
Audience	Teams that have completed an initial training and want continuity without turning AI into a separate project.

Learner outcomes

- Create a reliable place for practical questions and safe experiments.
- Reinforce good review habits as tools, interfaces, and model capabilities change.
- Turn scattered individual experiments into shared team learning.
- Maintain momentum without forcing AI into every workflow or chasing every announcement.

Guiding questions

- What did people try this month, and what happened?
- Which tool changes matter for this team, and which can be ignored for now?
- Where are review habits holding up, and where are outputs being trusted too quickly?
- What one practice should the team try before the next office hours session?

Readiness checks

- Participants bring real but non-sensitive questions and examples.
- The team can separate useful tool changes from noise.
- Repeated questions are converted into shared guidance or workflow notes.
- Office hours reinforce boundaries rather than encouraging careless experimentation.

Research-backed adoption practices

- Use recurring office hours as a learning system: collect questions, compare examples, update norms, and share what changed.
- Reward useful observations about failures and confusing outputs, not only successful prompts.
- Translate tool changes into work implications so people are not forced to chase every announcement.
- Revisit boundaries as tools gain new file, memory, browsing, coding, or agent capabilities.

Curriculum modules

Question intake

Each session starts with real questions, confusing moments, and examples from the previous month.

Tool-change translation

New AI features are explained in plain language, with emphasis on what is actually useful for the group.

Live workflow coaching

Participants work through non-sensitive scenarios and see how a workflow can be improved step by step.

Habit reinforcement

Each month ends with one or two practical habits the team can try before the next session.

Learning capture

Useful examples, cautions, and workflow refinements are captured so the team builds institutional memory over time.

Session flow

Collect

Gather questions, examples, confusing moments, and useful experiments from the previous month.

Translate

Explain relevant tool changes in practical language and ignore updates that do not affect the group.

Coach

Work through live non-sensitive examples, improving prompts and review habits together.

Capture

Document useful patterns, cautions, and candidate workflows.

Assign

Choose one small practice experiment for the next month.

Practice labs

Question intake board

Build: participant questions sorted by task, data boundary, tool class, review need, and next experiment. Use: real questions rewritten safely. Do: sort and turn two into practice prompts. Check: office hours stays practical and boundary-aware.

Feature-to-work triage

Build: a one-page digest translating one AI feature into who should care, what to test, and what to ignore. Use: a public product announcement. Do: map to tasks and design one safe test. Check: practice changes only when useful.

Failed-prompt rescue

Build: diagnosis sheet with failure category, missing context, revised prompt, review step, and reuse-or-stop decision. Use: a non-sensitive failed prompt or fictionalized version. Do: classify and revise. Check: the lesson explains the failure.

Monthly habit scorecard

Build: a scorecard tracking one habit, one safe experiment, one useful output, one failure, and one norm to update. Use: sanitized team practice notes. Do: apply the habit live and capture learning. Check: institutional memory grows.

1-2 hour utility projects

These short projects are designed to help individuals experience practical AI utility while keeping inputs safe and human review visible.

How to choose a strong first project

- Real work, safe input: choose a task learners recognize, then substitute public, fictional, or sanitized material.
- Named walk-away artifact: produce a file, table, checklist, brief, agenda, glossary, recipe card, or other object.
- Visible human judgment: require the learner to check claims, reject weak output, revise tone, name uncertainty, or decide what stays human-led.
- Transferable next use: teach a repeatable pattern such as adding context, requesting structure, comparing alternatives, verifying claims, documenting the recipe, and trying it again.

Public Article Learning Kit (60-75 minutes)

Turn one public article into a tutor session and a usable study packet.

Walk-away artifact: A one-page explainer, eight-term glossary, five-question quiz, misconception list, and verification table.

Safe input: Use a public article, help page, announcement, or topic summary. Avoid private client, patient, employee, financial, legal, or proprietary material.

- Paste or summarize one public source and state your current knowledge level.
- Ask for an explainer, glossary, examples, and common misconceptions.
- Create a five-question quiz and answer it before looking at the answer key.
- Mark every claim that needs source checking before you share or rely on it.

Starter prompt: Use this public source to create a learning kit for a smart beginner: [public source or topic]. Produce a one-page explainer, eight-term glossary, three examples, common misconceptions, a five-question quiz with answer key, and a verification table listing claims I should check in reliable sources.

Decision Brief in a Box (75-90 minutes)

Use AI to structure a decision without handing the decision away.

Walk-away artifact: A one-page decision brief with criteria, option table, assumption register, pre-mortem, and next-evidence checklist.

Safe input: Use a personal, public, fictional, or sanitized decision. Keep confidential finances, contracts, medical, legal, or customer details out of public tools.

- Name two or three options, the decision deadline, constraints, and what success would look like.
- Ask for a weighted criteria table and force the model to show its assumptions.
- Run a pre-mortem for the most tempting option and a second pass for the safest option.
- Write three human-only judgment questions you will answer before acting.

Starter prompt: Help me build a decision brief without deciding for me: [safe decision context]. Include weighted criteria, an option comparison table, assumptions, likely failure modes, missing evidence, who should be consulted, and three final human judgment questions.

Meeting-to-Momentum Kit (60-90 minutes)

Convert a rough meeting idea into a focused agenda and follow-through packet.

Walk-away artifact: A timed agenda, participant prep note, decision log template, action-item table, and follow-up email draft.

Safe input: Use fictional or sanitized meeting context. Do not paste private attendee notes, personnel details, patient data, customer records, or proprietary plans into public tools.

- Describe the meeting purpose, attendees by role, time available, and desired output.
- Ask for a timed agenda with decisions, discussion prompts, and parking-lot items.
- Create a follow-up template with owners, due dates, unresolved questions, and confirmation language.
- Delete or rewrite anything that sounds like a promise, personnel judgment, or confidential detail.

Starter prompt: Using this fictional or sanitized meeting context, create a meeting kit: timed agenda, prep questions by role, likely tensions, decision log, action-item table, and follow-up email draft. Mark assumptions and anything a human should confirm before sending: [context].

Before-and-After Writing Studio (60 minutes)

Use AI as an editor, not a ghostwriter, and leave with a cleaner real draft.

Walk-away artifact: A revised email, memo, bio, announcement, or article with a change log, tone rationale, and final review checklist.

Safe input: Use a rough draft that contains no sensitive personal, client, patient, employee, legal, financial, or proprietary details.

- Paste a safe rough draft and name the audience, purpose, tone, and length limit.
- Ask for two alternative revisions: one clearer and one warmer or more concise.
- Have AI critique the revision for unsupported claims, missing context, and awkward tone.
- Make the final edits yourself and keep a before-and-after note on what improved.

Starter prompt: Revise this safe rough draft for [audience] with a [tone] tone and a [length] limit. Give me two versions, explain the changes, then critique the stronger version for clarity, unsupported claims, missing context, and anything I should verify before sending: [safe draft].

Research Launch Board (90-120 minutes)

Plan the research before the rabbit hole opens.

Walk-away artifact: A research board with question clusters, search strings, source ladder, opposing views, red flags, and claim-check table.

Safe input: Use a public topic or general business question. Do not ask AI to invent citations or make final claims without sources.

- Start with one broad public question and ask AI to break it into subquestions.
- Generate search strings, source categories, keywords to avoid, and missing perspectives.
- Ask for a claim-check table with columns for claim, source needed, confidence, and next action.
- Use the board to search actual sources; do not treat the model output as evidence.

Starter prompt: Help me build a research launch board for [public topic]. Create question clusters, search strings, source types, likely blind spots, opposing perspectives, red flags, and a claim-check table. Do not invent citations or make final claims without sources.

Workflow Recipe Card (90-120 minutes)

Turn one repeated task into a reusable recipe your future self can follow.

Walk-away artifact: A recipe card with trigger, safe inputs, prompt sequence, review gates, owner, time saved estimate, and stop signs.

Safe input: Choose a recurring task and describe it generally or with fictional examples. Keep sensitive operational data out of public tools.

- Write the current task steps from trigger to final output.
- Mark which inputs are safe, which must be generalized, and which should not enter a public tool.
- Ask AI to propose a prompt sequence and review checkpoints for the safest useful slice.
- Test the recipe with one fictional example and revise the steps until they are repeatable.

Starter prompt: Turn this recurring task into a safe AI-assisted workflow recipe: [task]. Include trigger, user, safe inputs, prompt sequence, expected output, review checkpoints, human owner, estimated time saved, risks, stop signs, and when not to use AI.

Professional Proof Portfolio (75-90 minutes)

Use AI to turn scattered experience into concrete positioning you can defend.

Walk-away artifact: A 150-word bio, six resume bullets, proof bank, interview answer outlines, and claims-to-evidence checklist.

Safe input: Use a sanitized career summary or fictionalized role description. Do not include private employer data, references, compensation, or confidential projects.

- Write a sanitized inventory of roles, strengths, projects, and outcomes you can discuss publicly.
- Ask AI to draft positioning options and separate claims from evidence.
- Create resume bullets and interview outlines, then remove anything inflated or unverifiable.
- Practice three interview questions and revise the language until it sounds true in your own voice.

Starter prompt: Using this sanitized professional summary, help me build a proof portfolio. Draft a 150-word bio, six resume bullet options, a proof bank, three interview answer outlines, and a claims-to-evidence checklist. Flag anything that sounds inflated or needs stronger evidence: [summary].

Seven-Day Action Plan (60-75 minutes)

Turn a vague goal into a week of small, scheduled actions.

Walk-away artifact: A seven-day plan with calendar blocks, next actions, obstacle responses, accountability note, and end-of-week review.

Safe input: Use personal goals and constraints that are not private, medical, financial, legal, or deeply sensitive.

- Describe the goal, available time, constraints, and what usually derails progress.
- Ask for a one-week plan with daily actions no longer than 45 minutes.
- Have AI identify likely obstacles and create if-then responses.
- Put the first two actions on your calendar and delete anything unrealistic.

Starter prompt: Help me turn this safe goal into a seven-day action plan: [safe goal]. Ask up to five clarifying questions if needed, then create calendar blocks, daily next actions under 45 minutes, obstacle responses, an accountability note, and an end-of-week review checklist.

Tangible cases

Prompt that failed

Situation: A team member tried AI for a harmless task, but the result was generic, wrong, or too confident.

Learner task: Create a failed-prompt diagnosis sheet, revised prompt, review checkpoint, and reuse-or-stop decision.

Starter prompt: Diagnose why this AI attempt failed. Create a table with failure category, evidence from the output, missing context, revised prompt, review checkpoint, and reuse-or-stop recommendation. Categories include task choice, missing context, unclear constraints, weak examples, overtrust, and review failure: [failed prompt and output].

New feature translation

Situation: A model or product announces a new feature, and the team is unsure whether it matters.

Learner task: Create a feature-to-work digest naming who should care, what to test, what to ignore, and what boundary changed.

Starter prompt: Explain this AI tool update for a cautious team. Produce a one-page digest with what changed, who should care, likely work impact, what to ignore, remaining risks or boundaries, and one small safe experiment to try: [update].

Shared learning capture

Situation: Several staff members are experimenting privately, but the organization is not learning from those experiments.

Learner task: Convert individual examples into a team recap, reusable patterns, caution list, workflow candidates, and next questions.

Starter prompt: Turn these monthly AI practice notes into a team learning recap. Include useful patterns, confusing moments, unsafe ideas to avoid, reusable prompts, workflow candidates, owner questions, and a next-month practice habit: [practice notes].

Prompt library

Question refinement

Turn this messy AI question into a clearer office-hours question. Identify the real task, safe substitute input, missing context, data boundaries, desired artifact, review need, and whether it belongs in live demo, policy owner, or later research: [question].

Workflow rescue

Rescue this disappointing AI workflow. Identify where it broke down, rewrite the prompt sequence, add safe-input guidance, add review checkpoints, define the desired artifact, and say when the task should stay human-led: [workflow].

Tool-change digest

Summarize this AI tool change for a practical team. Use five sections: what changed, who should care, what work it may affect, what to ignore for now, and one safe test for next month: [announcement].

Monthly recap

Create a one-page office-hours recap from these notes. Include questions answered, artifacts created, examples discussed, practice habit, cautions, owner follow-ups, and items to revisit next month: [notes].

Materials to develop

- Monthly question intake form
- Office-hours recap template
- Tool-change digest format
- Team practice log
- Workflow refinement log
- Reusable question bank

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- Monthly learning loop
- Question-to-practice coaching flow
- Office-hours knowledge capture cycle

Follow-up work

- Try the monthly practice habit with one safe task.
- Submit questions or examples before the next session.
- Update shared notes with what worked, what failed, and what needs clearer guidance.

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.