

AI Fluency Essentials

A calm, practical introduction for people who want to understand LLMs, use them safely, and build confidence with ordinary work tasks.

Format	One 90-minute Zoom session with guided practice and a follow-up handout.
Level	Level 1: shared team AI fluency
Audience	Curious professionals, mixed-comfort teams, and organizations that need a shared starting point before deeper tool decisions.

Learner outcomes

- Explain what LLMs are good at, where they fail, and why human review matters.
- Use a basic prompting loop for drafting, summarizing, planning, and asking better questions.
- Recognize sensitive-data boundaries and safer ways to practice with non-sensitive examples.
- Leave with shared vocabulary for prompts, context, output quality, uncertainty, and review.

Guiding questions

- What does our team need to understand before choosing tools or policies?
- Which everyday tasks are safe places to build skill?
- How should we review AI output before it affects real work?
- What shared language will help us discuss AI without hype or fear?

Readiness checks

- Participants can describe AI as an assistive drafting and reasoning partner, not an authority.
- Participants can use a repeatable prompt structure.
- Participants can explain the difference between safe examples and sensitive operational data.
- Participants can name a human review habit they will use before applying output.

Curriculum modules

The practical mental model

Participants learn how prompts, context, instructions, examples, and review shape an AI session without needing technical background.

Prompting as conversation design

We practice turning vague requests into useful context, constraints, examples, and review questions.

Safe experimentation

The session establishes clear boundaries around patient, customer, employee, legal, financial, and proprietary information.

Confidence routines

Participants leave with small repeatable habits: draft, critique, revise, verify, and decide what still needs human judgment.

Team language and norms

The group names simple shared expectations for AI use, including what to try, what to avoid, and how to talk about uncertain outputs.

Session flow

Baseline

Surface current comfort levels, hopes, concerns, and common misconceptions.

Model

Teach a practical mental model for context, instructions, examples, model output, and review.

Practice

Run guided prompt exercises using neutral scenarios that resemble the team's work without exposing sensitive data.

Review

Apply a quality checklist to outputs and discuss what should be verified, revised, or discarded.

Normalize

Close with team language, safe-use reminders, and a first set of practice habits.

Practice labs

Prompt anatomy lab

Rewrite a weak prompt into a stronger one by adding role, task, context, audience, constraints, and success criteria.

Output comparison

Compare two model responses and identify which is more useful, what is missing, and what should be checked.

Data-boundary sorting

Sort example inputs into safe public practice, caution, and private-data categories.

Revision ladder

Practice moving from first draft to better draft through critique, constraint, and audience-shift prompts.

Tangible cases

Shared vocabulary kickoff

Situation: A leadership group and frontline staff are using different words for prompts, outputs, hallucinations, and review.

Learner task: Create a shared glossary and test it against everyday examples so the group can discuss AI without confusion.

Starter prompt: Create a plain-language glossary for a team learning AI. Define prompt, context, output, hallucination, review, sensitive data, and human judgment. For each term, give a workplace example and a common misconception.

Nonprofit communications draft

Situation: A nonprofit team wants to draft a donor update from public program facts and a fictional event recap.

Learner task: Practice drafting, audience adaptation, and review without using donor records or private beneficiary details.

Starter prompt: Draft a donor update using only the public facts and fictional event notes below. Make it warm, specific, and modest. Then list what a staff member should verify before sending: [public facts and fictional notes].

Healthcare admin FAQ practice

Situation: A cautious health care admin team wants AI practice that does not involve patient information.

Learner task: Use a fictional policy FAQ to practice summarizing, tone adjustment, and boundary setting.

Starter prompt: Using this fictional clinic FAQ, rewrite the answer for a general audience. Do not add medical advice. Flag anything that should be reviewed by the appropriate internal owner before publication: [fictional FAQ].

Prompt library

Team baseline

Ask five questions that help a mixed-comfort team describe what they know about AI, what worries them, and what they want to learn. Keep the questions nontechnical and practical.

Prompt anatomy

Rewrite this vague prompt into a stronger one with role, task, audience, context, constraints, examples, and success criteria. Explain what changed and why: [rough prompt].

Output comparison

Compare these two AI outputs. Which is more useful, more accurate, clearer, and safer to use? Name what still needs human review before either one is applied: [output A] [output B].

Boundary sort

Sort these example inputs into safe public practice, caution, and do not paste into public AI tools. Give a short reason for each decision: [example inputs].

Materials to develop

- AI fluency starter guide
- Prompt pattern quick sheet
- Sensitive-data boundary checklist
- Output review checklist
- Team language and norms worksheet

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- LLM session loop: prompt, context, output, review, revision
- Safe-data boundary map for public AI tools
- Team fluency ladder from curiosity to useful practice

Follow-up work

- Run one safe prompt experiment before the next meeting.
- Collect examples of confusing, useful, or unreliable AI outputs for discussion.
- Identify one team workflow that may deserve deeper workshop treatment.

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.