

Clear Prompting Field Guide

Better prompting is not a secret vocabulary or a clever trick. It is clearer thinking in sentence form: context, vocabulary, logic, output shape, and review criteria. This guide is meant to feel practical on the first read. You should be able to copy a prompt, adapt it to safe material, and immediately see why the result is easier to review.

The deeper habit is simple: do not ask AI to guess the job. Tell it the job, the reader, the constraints, the shape of the answer, and how the answer will be checked. That small amount of up-front clarity saves time because it reduces retry loops, and it prevents blunders because it makes assumptions visible before output is used.

Best use	A friendly first handout for AI-curious individuals, teams, and leaders who want practical wins without hype.
Time payoff	Fewer vague retries, faster review, and prompt cards that can be reused in real workflows.
Blunders avoided	Invented facts, wrong audience, hidden assumptions, shifted meaning, accidental promises, and outputs that sound finished before they are reviewed.

How to use this guide

- Read the six moves first. They are the habits behind nearly every good prompt.
- Skim the poor-vs-excellent examples to see how a vague request becomes a usable work instruction.
- Copy one ready prompt, replace the fictional context with safe material, and run it as a low-risk practice exercise.
- Review the output before using it. Clear prompting improves the draft, but it does not remove human responsibility.
- Save any prompt that survives review. A useful prompt is a small workflow asset.

The premise

- The model cannot know which version of a task you mean unless you name the context.
- The words that seem obvious to humans often carry hidden standards that should be defined.
- Small logic words such as only, before, unless, preserve, compare, and verify change the behavior of the output.

- The shape of the answer is part of the work: table, checklist, brief, revision log, or decision memo.
- A strong prompt includes its own review habit so the user does not mistake fluent output for finished work.

Six clear prompting moves

Context narrows the universe

Tell the model the situation, audience, purpose, constraints, prior attempts, and what the output will be used for. Without context, the model answers an average version of the request. With context, it can aim at the actual work in front of you.

Payoff: Less generic output, fewer rounds of repair, and a result that starts closer to the real work.

Vocabulary carries judgment

Define words such as brief, warm, rigorous, practical, leadership-ready, or beginner-friendly when those standards matter. Many prompt failures happen because a familiar word quietly carries a standard the model cannot infer.

Payoff: The model stops guessing your standards and starts working inside them.

Sentence logic controls behavior

Use precise relationships: compare this with that, do X before Y, only use these facts, ask questions if evidence is missing, and do not decide for me. Small logic words such as only, before, unless, preserve, avoid, and verify are not decorations. They steer the task.

Payoff: You prevent hidden assumptions, invented commitments, and answers that solve the wrong problem.

Output shape saves time

Ask for a table, checklist, memo, decision brief, revision log, or ordered plan when that is what the work needs. A paragraph may sound polished, but a table or checklist is often easier to inspect, assign, revise, and reuse.

Payoff: The answer becomes easier to scan, edit, assign, verify, and reuse.

Review criteria reduce blunders

Name how the output will be judged: accuracy risk, tone, missing context, audience fit, unsupported claims, data boundaries, and human approval needs. A useful prompt can produce the draft and the checklist for inspecting the draft.

Payoff: The same prompt that generates work can also make the work safer to inspect.

Iteration turns effort into assets

Save the version that worked, name when it should be used, and note what still requires human judgment. Clear prompting becomes valuable when a useful experiment turns into a reusable habit, not when a single answer sounds impressive.

Payoff: A single experiment becomes a reusable prompt card, workflow recipe, or specialty assistant seed.

Poor prompts and excellent prompts

The goal is not to memorize these exact prompts. The goal is to notice how context, vocabulary, sentence logic, output shape, and review criteria change the quality of the work.

Business email revision

Situation: A professional needs to rewrite a sensitive but non-confidential email without sounding vague, cold, or overpromising.

Poor: Make this email better.

Excellent: Revise this fictional email for a busy department director. Keep the message under 160 words, make the tone calm and direct, preserve the three facts listed below, avoid promising a deadline, and return two versions: concise and warmer. Then list what a human should confirm before sending. Facts to preserve: [facts]. Draft: [email].

Why it works: Names audience, tone, length, fixed facts, a promise to avoid, output format, and review criteria.

Blunder avoided: Prevents invented commitments or polished language that changes the facts.

Time gain: Cuts several vague revision rounds down to one focused comparison and a review checklist.

Article learning kit

Situation: A learner wants to understand a public article quickly without mistaking a fluent summary for verified knowledge.

Poor: Summarize this article.

Excellent: Create a learning kit from this public article for a smart beginner. Include a 150-word summary, eight-term glossary, three practical examples, five-question quiz with answer key, common misconceptions, and a claim-check table. Separate what the article says from what would need outside verification. Article: [public article].

Why it works: Turns summarization into a structured learning task with vocabulary, examples, assessment, and verification.

Blunder avoided: Prevents passive acceptance of confident claims by making source checking visible.

Time gain: Converts reading support into a reusable study packet in one pass.

Meeting follow-through

Situation: A team wants momentum after a meeting but has only rough, sanitized notes.

Poor: Turn these notes into action items.

Excellent: Using these fictional meeting notes, create a follow-through packet with five sections: decision summary, unresolved questions, action-item table with owner and deadline placeholders, risks or assumptions to confirm, and a follow-up email draft. Do not invent owners, dates, or decisions. Mark anything that needs human confirmation. Notes: [notes].

Why it works: Defines the deliverable, separates decisions from open questions, and blocks invented commitments.

Blunder avoided: Prevents false certainty in action items, a common source of team confusion.

Time gain: Turns messy notes into a scannable packet the team can confirm quickly.

Research planning

Situation: A leader wants to explore an unfamiliar topic before reading sources or talking to vendors.

Poor: Research AI tools for us.

Excellent: Help me plan research on AI tools for a cautious midsize organization. Do not recommend products yet. Create question clusters, search terms, source categories, comparison criteria, risk questions, vendor-claim questions, and a table of claims we must verify in current sources. End with the first five searches a human should run.

Why it works: Keeps the model in planning mode and prevents premature recommendations.

Blunder avoided: Avoids tool-chasing before the organization understands its own use cases.

Time gain: Turns a broad research fog into a search plan and comparison structure in minutes.

Decision support

Situation: A person or team needs help thinking through options without handing the decision to AI.

Poor: Which option should we choose?

Excellent: Help us think through these options without choosing for us. Create a decision brief with criteria, option comparison table, assumptions, missing evidence, likely failure modes, stakeholder questions, and final human decision questions. If facts are missing, ask for them or label the gap. Options: [options].

Why it works: Uses AI for structure, comparison, and blind-spot detection while keeping accountability human.

Blunder avoided: Prevents outsourcing judgment to a fluent answer that lacks context or responsibility.

Time gain: Compresses the first draft of a decision memo while making remaining human work clearer.

Actual prompts you can copy, edit, and reuse

These prompts are intentionally concrete. They use fictional or public-safe contexts so learners can practice without exposing patient, customer, employee, financial, legal, or proprietary material. Copy one, change the safe context, and keep the review instructions intact.

Warm follow-up email

Use when: When you need a polished message but must preserve a few facts and avoid accidental promises.

Prompt: Revise this fictional email for a warm, concise follow-up to a professional who missed an AI fluency workshop. Audience: busy adult learner. Purpose: acknowledge the missed session, offer useful next steps, avoid guilt, and do not promise a recording. Facts to preserve: no recording is available; a one-page notes sheet can be sent; the next office-hours session is Thursday at 2:00 p.m.; questions are welcome by email. Return two versions under 140 words and a send-before-review checklist.

Why it works: It names audience, tone, purpose, fixed facts, a promise to avoid, length, output count, and review needs.

Meeting-to-action packet

Use when: When a meeting needs momentum but the notes are rough and should not be treated as confirmed truth.

Prompt: Use this fictional meeting context to create a follow-through packet. Context: a small nonprofit team discussed a fall volunteer orientation, a donor update, and a website FAQ refresh. Create five sections: decision summary, unresolved questions, action-item table with owner and date placeholders, risks or assumptions to confirm, and a follow-up email draft. Do not invent owners, dates, budgets, or final decisions. Mark every item that needs human confirmation.

Why it works: It turns messy notes into a useful artifact while blocking false certainty about owners, dates, and decisions.

Public article learning kit

Use when: When a learner wants to understand a public source without mistaking a fluent summary for verified knowledge.

Prompt: Create a learning kit from a public article about urban heat islands for a smart beginner. Include a 150-word summary, eight-term glossary, three practical examples, five-question quiz with answer key, common misconceptions, and a claim-check table. Separate what the article says from what would need outside verification. Use only the article text I provide, and label any point that requires another source.

Why it works: It asks for learning support, practice questions, and verification boundaries instead of a shapeless summary.

Decision brief without outsourcing judgment

Use when: When a person or team needs help thinking clearly but must keep the decision human-owned.

Prompt: Help me think through three options without choosing for me: a one-time team briefing, monthly AI office hours, or waiting three months. Create a decision brief with criteria, comparison table, assumptions,

missing evidence, likely failure modes, stakeholder questions, and final human decision questions. If a fact is missing, label the gap instead of filling it in.

Why it works: It uses AI for structure, comparison, and blind-spot detection while refusing to hand over the decision.

Plain-language FAQ rewrite

Use when: When cautious organizations need clearer writing without legal, medical, or compliance overreach.

Prompt: Rewrite this fictional clinic admin FAQ for a general audience. Keep the meaning intact, use plain language, create a jargon glossary, list questions that require an internal owner, and flag phrases that could sound like medical, legal, privacy, security, or compliance advice. Do not add new policy claims. End with a review checklist for the appropriate staff owner.

Why it works: It makes clarity useful while naming the boundaries that require human and organizational review.

Prompt repair coach

Use when: When a learner has a vague request and needs help turning it into a reliable prompt card.

Prompt: Turn this rough request into a clearer prompt card: "Help me make our onboarding better." Ask up to five clarifying questions first if needed. Then produce a prompt with job, audience, context needed, key vocabulary to define, constraints, output format, review criteria, and what the model must not do. Finish by explaining why each part improves the prompt.

Why it works: It teaches the structure behind better prompting instead of merely producing one improved sentence.

Prompting tips that repay attention

- Name the job before naming the tool: explain, compare, revise, critique, plan, extract, translate, summarize, or decide what to ask next.
- Give the model a reader: board member, patient-facing staff member, donor, technician, beginner, or skeptical manager.
- Define the words that carry standards. If clear means brief or rigorous means evidence-tagged, say so.
- Protect facts that must not move by listing them separately from material that may be rewritten.
- Use complete constraints: under 180 words, no promises, calm closing sentence, and no new facts is clearer than short and professional.
- Separate creating from judging. Ask for a draft first, then ask for a critique.
- Ask for questions when context is missing instead of letting the model fill the gaps silently.
- Choose the output container: table, checklist, brief, log, comparison, or revision memo.
- Make uncertainty visible by asking for assumptions, missing facts, confidence level, and claims to verify.
- Keep private data out of public tools unless an approved private environment and internal owner are in place.
- Use realistic but safe examples. A fictional donor update or sanitized meeting note teaches more than an abstract toy prompt.
- Request a change log when revising so accidental meaning shifts are easier to catch.
- Ask the model to explain its own interpretation of your request when the stakes are higher or the task is easy to misunderstand.
- Use the first answer as raw material. The goal is not to accept the first draft, but to get something structured enough to review.
- Save prompts that survive review. Name when to use them and what a human still needs to check.

Practice exercise: repair a vague prompt

Choose one poor prompt you have used recently. Do not start by making it fancy. Start by making it inspectable. Rewrite it with job, audience, context, vocabulary definitions, constraints, output format, review criteria, and a sentence saying what the model must not do.

- Step 1: Write the rough request exactly as you might normally type it.
- Step 2: Add the real reader, purpose, and safe background context.
- Step 3: Define two or three words that carry standards, such as concise, warm, rigorous, practical, or executive-ready.
- Step 4: Choose the output shape: table, checklist, memo, brief, lesson plan, revision log, or action list.

- Step 5: Add review instructions: what to verify, what to avoid, what assumptions to label, and what remains human-owned.
- Step 6: Save the improved version only if the output survives review.

Starter prompt for the exercise

Turn my rough request into a clearer prompt. Ask up to five clarifying questions if needed. Then produce a prompt card with job, audience, context, key vocabulary, constraints, output format, review criteria, and what not to do. Rough request: [paste request].

Contact

For guided AI fluency training over Zoom: contact@xensible.com