

Executive AI Briefing

Plain-spoken AI orientation for leaders who need clarity before making budget, policy, vendor, or training decisions.

Format	One focused Zoom briefing, optionally followed by a Q&A; or planning session.
Level	Leadership level: clarity before technology decisions
Audience	Executives, founders, board-adjacent leaders, and department heads who need a useful map rather than hype.

Learner outcomes

- Separate AI training needs from software purchasing needs.
- Ask better questions about vendors, data, policy, implementation, and organizational readiness.
- Identify decisions that need more learning before investment.
- Leave with a leadership vocabulary for opportunity, uncertainty, human review, and responsible experimentation.

Guiding questions

- What should leaders understand before authorizing tools, pilots, policies, or budgets?
- Where does education need to come before procurement?
- Which risks require the right internal owners rather than a training promise?
- What small next step would increase clarity without locking the organization into a premature decision?

Readiness checks

- Leaders can explain why fluency training may be needed before tool investment.
- Leaders can ask practical questions about data boundaries, review, ownership, and expected value.
- The organization has named which decisions require internal policy, legal, privacy, or technical owners.
- The next step is small enough to learn from and clear enough to communicate.

Research-backed adoption practices

- Treat AI adoption as work redesign and capability building, not simply tool access.
- Ask what evidence would justify scaling before approving major spend or vendor lock-in.
- Align incentives, manager support, and evaluation habits so experimentation can produce shared learning.
- Create a basic evaluation structure: who reviews outputs, who updates workflows, and how lessons are captured.

Curriculum modules

What leaders need to know now

A clear map of current AI capabilities, limitations, tool categories, and common misconceptions.

Clarity before complexity

We distinguish fluency, workflow practice, policy, vendor evaluation, and implementation work.

Responsible decision questions

Leaders practice asking about data exposure, review requirements, accountability, and expected business value.

Next-step decision paths

The briefing ends with a recommendation for training, discovery, a small pilot, or a deliberate wait.

Internal communication

Leaders practice explaining AI next steps in language that reduces confusion, fear, and unrealistic expectations.

Session flow

Landscape

Explain current AI capabilities, limitations, and tool categories in plain language.

Separate

Distinguish fluency training, workflow practice, policy, vendor evaluation, implementation, and governance.

Question

Work through leader-level questions about value, data exposure, review burden, and accountability.

Decide

Sort likely next steps into train, discover, pilot, wait, or seek specialist advice.

Communicate

Draft a short internal message that frames AI learning without hype or guarantees.

Practice labs

Vendor claim scorecard

Build: a scorecard turning one vendor claim into workflow-fit, evidence, data, review, and implementation questions. Use: a public claim or fictional pitch. Do: extract implied promises and evidence requests. Check: no legal, security, or compliance conclusions.

Decision-gate worksheet

Build: a train, discover, pilot, buy, wait, or specialist-review worksheet for one AI idea. Use: a fictional or sanitized idea. Do: separate learning, workflow, tool, and governance needs. Check: smallest next step avoids lock-in.

Risk language rewrite

Build: before-and-after leadership language naming uncertainty, boundaries, and owner responsibilities. Use: a fictional or sanitized AI message. Do: remove hype and vague risk language. Check: reassurance comes from specificity, not guarantees.

Leadership message draft

Build: a 250-word internal message explaining why the organization is learning first, what is in bounds, and what comes next. Use: sanitized organizational context. Do: draft and revise for warmth and clarity. Check: anxiety and hype both go down.

1-2 hour utility projects

These short projects are designed to help individuals experience practical AI utility while keeping inputs safe and human review visible.

How to choose a strong first project

- Real work, safe input: choose a task learners recognize, then substitute public, fictional, or sanitized material.
- Named walk-away artifact: produce a file, table, checklist, brief, agenda, glossary, recipe card, or other object.
- Visible human judgment: require the learner to check claims, reject weak output, revise tone, name uncertainty, or decide what stays human-led.
- Transferable next use: teach a repeatable pattern such as adding context, requesting structure, comparing alternatives, verifying claims, documenting the recipe, and trying it again.

Public Article Learning Kit (60-75 minutes)

Turn one public article into a tutor session and a usable study packet.

Walk-away artifact: A one-page explainer, eight-term glossary, five-question quiz, misconception list, and verification table.

Safe input: Use a public article, help page, announcement, or topic summary. Avoid private client, patient, employee, financial, legal, or proprietary material.

- Paste or summarize one public source and state your current knowledge level.
- Ask for an explainer, glossary, examples, and common misconceptions.
- Create a five-question quiz and answer it before looking at the answer key.
- Mark every claim that needs source checking before you share or rely on it.

Starter prompt: Use this public source to create a learning kit for a smart beginner: [public source or topic]. Produce a one-page explainer, eight-term glossary, three examples, common misconceptions, a five-question quiz with answer key, and a verification table listing claims I should check in reliable sources.

Decision Brief in a Box (75-90 minutes)

Use AI to structure a decision without handing the decision away.

Walk-away artifact: A one-page decision brief with criteria, option table, assumption register, pre-mortem, and next-evidence checklist.

Safe input: Use a personal, public, fictional, or sanitized decision. Keep confidential finances, contracts, medical, legal, or customer details out of public tools.

- Name two or three options, the decision deadline, constraints, and what success would look like.
- Ask for a weighted criteria table and force the model to show its assumptions.
- Run a pre-mortem for the most tempting option and a second pass for the safest option.
- Write three human-only judgment questions you will answer before acting.

Starter prompt: Help me build a decision brief without deciding for me: [safe decision context]. Include weighted criteria, an option comparison table, assumptions, likely failure modes, missing evidence, who should be consulted, and three final human judgment questions.

Meeting-to-Momentum Kit (60-90 minutes)

Convert a rough meeting idea into a focused agenda and follow-through packet.

Walk-away artifact: A timed agenda, participant prep note, decision log template, action-item table, and follow-up email draft.

Safe input: Use fictional or sanitized meeting context. Do not paste private attendee notes, personnel details, patient data, customer records, or proprietary plans into public tools.

- Describe the meeting purpose, attendees by role, time available, and desired output.
- Ask for a timed agenda with decisions, discussion prompts, and parking-lot items.
- Create a follow-up template with owners, due dates, unresolved questions, and confirmation language.

- Delete or rewrite anything that sounds like a promise, personnel judgment, or confidential detail.

Starter prompt: Using this fictional or sanitized meeting context, create a meeting kit: timed agenda, prep questions by role, likely tensions, decision log, action-item table, and follow-up email draft. Mark assumptions and anything a human should confirm before sending: [context].

Before-and-After Writing Studio (60 minutes)

Use AI as an editor, not a ghostwriter, and leave with a cleaner real draft.

Walk-away artifact: A revised email, memo, bio, announcement, or article with a change log, tone rationale, and final review checklist.

Safe input: Use a rough draft that contains no sensitive personal, client, patient, employee, legal, financial, or proprietary details.

- Paste a safe rough draft and name the audience, purpose, tone, and length limit.
- Ask for two alternative revisions: one clearer and one warmer or more concise.
- Have AI critique the revision for unsupported claims, missing context, and awkward tone.
- Make the final edits yourself and keep a before-and-after note on what improved.

Starter prompt: Revise this safe rough draft for [audience] with a [tone] tone and a [length] limit. Give me two versions, explain the changes, then critique the stronger version for clarity, unsupported claims, missing context, and anything I should verify before sending: [safe draft].

Research Launch Board (90-120 minutes)

Plan the research before the rabbit hole opens.

Walk-away artifact: A research board with question clusters, search strings, source ladder, opposing views, red flags, and claim-check table.

Safe input: Use a public topic or general business question. Do not ask AI to invent citations or make final claims without sources.

- Start with one broad public question and ask AI to break it into subquestions.
- Generate search strings, source categories, keywords to avoid, and missing perspectives.
- Ask for a claim-check table with columns for claim, source needed, confidence, and next action.
- Use the board to search actual sources; do not treat the model output as evidence.

Starter prompt: Help me build a research launch board for [public topic]. Create question clusters, search strings, source types, likely blind spots, opposing perspectives, red flags, and a claim-check table. Do not invent citations or make final claims without sources.

Workflow Recipe Card (90-120 minutes)

Turn one repeated task into a reusable recipe your future self can follow.

Walk-away artifact: A recipe card with trigger, safe inputs, prompt sequence, review gates, owner, time saved estimate, and stop signs.

Safe input: Choose a recurring task and describe it generally or with fictional examples. Keep sensitive operational data out of public tools.

- Write the current task steps from trigger to final output.
- Mark which inputs are safe, which must be generalized, and which should not enter a public tool.
- Ask AI to propose a prompt sequence and review checkpoints for the safest useful slice.
- Test the recipe with one fictional example and revise the steps until they are repeatable.

Starter prompt: Turn this recurring task into a safe AI-assisted workflow recipe: [task]. Include trigger, user, safe inputs, prompt sequence, expected output, review checkpoints, human owner, estimated time saved, risks, stop signs, and when not to use AI.

Professional Proof Portfolio (75-90 minutes)

Use AI to turn scattered experience into concrete positioning you can defend.

Walk-away artifact: A 150-word bio, six resume bullets, proof bank, interview answer outlines, and claims-to-evidence checklist.

Safe input: Use a sanitized career summary or fictionalized role description. Do not include private employer data, references, compensation, or confidential projects.

- Write a sanitized inventory of roles, strengths, projects, and outcomes you can discuss publicly.
- Ask AI to draft positioning options and separate claims from evidence.
- Create resume bullets and interview outlines, then remove anything inflated or unverifiable.
- Practice three interview questions and revise the language until it sounds true in your own voice.

Starter prompt: Using this sanitized professional summary, help me build a proof portfolio. Draft a 150-word bio, six resume bullet options, a proof bank, three interview answer outlines, and a claims-to-evidence checklist. Flag anything that sounds inflated or needs stronger evidence: [summary].

Seven-Day Action Plan (60-75 minutes)

Turn a vague goal into a week of small, scheduled actions.

Walk-away artifact: A seven-day plan with calendar blocks, next actions, obstacle responses, accountability note, and end-of-week review.

Safe input: Use personal goals and constraints that are not private, medical, financial, legal, or deeply sensitive.

- Describe the goal, available time, constraints, and what usually derails progress.
- Ask for a one-week plan with daily actions no longer than 45 minutes.
- Have AI identify likely obstacles and create if-then responses.

- Put the first two actions on your calendar and delete anything unrealistic.

Starter prompt: Help me turn this safe goal into a seven-day action plan: [safe goal]. Ask up to five clarifying questions if needed, then create calendar blocks, daily next actions under 45 minutes, obstacle responses, an accountability note, and an end-of-week review checklist.

Tangible cases

Vendor pitch review

Situation: A CEO hears a vendor promise broad productivity gains and wants better questions before spending money.

Learner task: Translate the pitch into a vendor scorecard with evidence requests, workflow-fit questions, and training implications.

Starter prompt: Turn this vendor claim into a leadership scorecard. Include claimed value, evidence to request, workflow-fit questions, data-handling questions, human-review requirements, implementation burden, training needs, lock-in concerns, and what would justify a pilot: [vendor claim].

Board-level AI stance

Situation: A board or senior leadership group asks whether the organization has an AI strategy.

Learner task: Draft a board-ready learning stance with boundaries, next-step gates, and owner roles without overpromising outcomes.

Starter prompt: Draft a board-level AI learning statement of 250 words or fewer. Include clarity before complexity, staff fluency, safe experimentation, human judgment, owner roles, near-term training steps, and decisions that require specialist review. Avoid hype and guarantees.

Policy before practice tension

Situation: A department wants to ban, buy, or standardize AI before staff understand practical use.

Learner task: Create a decision table separating training, workflow practice, policy, privacy/legal, technical, procurement, and wait items.

Starter prompt: Given this leadership concern, create a decision table with columns for issue, category, owner, risk, evidence needed, and smallest next step. Categories: fluency training, workflow practice, policy, privacy/legal, technical implementation, procurement, specialist review, or wait: [concern].

Prompt library

Decision gate

Evaluate this AI idea through decision gates: learn, practice, discover, pilot, buy, wait, or seek specialist review. Produce a table with gate, evidence needed, owner, risk, cost of being wrong, and next action: [AI idea].

Leadership briefing prep

Prepare a plain-language executive briefing on [AI topic]. Include what it is, what it is useful for, where it fails, practical examples, limits, risks, questions leaders should ask, training implications, and a cautious next step.

Vendor question bank

Create a vendor question bank for [tool category]. Organize it by workflow fit, evidence, data use, privacy/security review, human review, admin controls, implementation effort, training, support, lock-in, and exit options.

Internal message

Draft an internal message explaining that we are learning about AI before making major technology decisions. Keep it calm, practical, nontechnical, and clear about responsible use boundaries, safe practice examples, what not to paste into public tools, and where questions should go.

Materials to develop

- Executive briefing deck outline
- Vendor question guide
- AI investment readiness checklist
- Train, pilot, buy, or wait worksheet
- Leadership communication template
- Decision-gate question bank

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- Decision gates for AI investment
- Build, buy, train, or wait decision tree
- Leadership clarity map: learn, test, decide, communicate
- Responsible adoption operating model: people, process, governance, measurement

Follow-up work

- Select a first audience for training or discovery.
- List vendor, policy, or data questions that require specialist review.
- Schedule a follow-up briefing or readiness sprint if the organization needs a fuller map.

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.