

AI Fluency Essentials

A calm, practical introduction for people who want to understand LLMs, use them safely, and build confidence with ordinary work tasks.

Format	One 90-minute Zoom session with guided practice and a follow-up handout.
Level	Level 1: shared team AI fluency
Audience	Curious professionals, mixed-comfort teams, and organizations that need a shared starting point before deeper tool decisions.

Learner outcomes

- Explain what LLMs are good at, where they fail, and why human review matters.
- Use a basic prompting loop for drafting, summarizing, planning, and asking better questions.
- Recognize sensitive-data boundaries and safer ways to practice with non-sensitive examples.
- Leave with shared vocabulary for prompts, context, output quality, uncertainty, and review.

Guiding questions

- What does our team need to understand before choosing tools or policies?
- Which everyday tasks are safe places to build skill?
- How should we review AI output before it affects real work?
- What shared language will help us discuss AI without hype or fear?

Readiness checks

- Participants can describe AI as an assistive drafting and reasoning partner, not an authority.
- Participants can use a repeatable prompt structure.
- Participants can explain the difference between safe examples and sensitive operational data.
- Participants can name a human review habit they will use before applying output.

Curriculum modules

The practical mental model

Participants learn how prompts, context, instructions, examples, and review shape an AI session without needing technical background.

Prompting as conversation design

We practice turning vague requests into useful context, constraints, examples, and review questions.

Safe experimentation

The session establishes clear boundaries around patient, customer, employee, legal, financial, and proprietary information.

Confidence routines

Participants leave with small repeatable habits: draft, critique, revise, verify, and decide what still needs human judgment.

Team language and norms

The group names simple shared expectations for AI use, including what to try, what to avoid, and how to talk about uncertain outputs.

Session flow

Baseline

Surface current comfort levels, hopes, concerns, and common misconceptions.

Model

Teach a practical mental model for context, instructions, examples, model output, and review.

Practice

Run guided prompt exercises using neutral scenarios that resemble the team's work without exposing sensitive data.

Review

Apply a quality checklist to outputs and discuss what should be verified, revised, or discarded.

Normalize

Close with team language, safe-use reminders, and a first set of practice habits.

Practice labs

Prompt anatomy lab

Rewrite a weak prompt into a stronger one by adding role, task, context, audience, constraints, and success criteria.

Output comparison

Compare two model responses and identify which is more useful, what is missing, and what should be checked.

Data-boundary sorting

Sort example inputs into safe public practice, caution, and private-data categories.

Revision ladder

Practice moving from first draft to better draft through critique, constraint, and audience-shift prompts.

Materials to develop

- AI fluency starter guide
- Prompt pattern quick sheet
- Sensitive-data boundary checklist
- Output review checklist
- Team language and norms worksheet

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- LLM session loop: prompt, context, output, review, revision
- Safe-data boundary map for public AI tools
- Team fluency ladder from curiosity to useful practice

Follow-up work

- Run one safe prompt experiment before the next meeting.
- Collect examples of confusing, useful, or unreliable AI outputs for discussion.
- Identify one team workflow that may deserve deeper workshop treatment.

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.