

Free AI Fluency Starter

A public starter path for people who want a clear, non-intimidating first step into LLMs, prompting, safe practice, and practical review habits.

Format	Self-guided webpage and downloadable PDF companion, designed as a gentle entry point before a team briefing or coached workshop.
Level	Level 0: orientation and first safe practice
Audience	Individuals, leaders, and teams who are curious about AI but want plain-language basics before committing to training.

Learner outcomes

- Understand what LLMs do, why they can be useful, and why their answers still need human review.
- Use a simple prompt structure for common tasks such as drafting, summarizing, planning, and asking better questions.
- Practice with non-sensitive examples while keeping patient, customer, employee, legal, financial, and proprietary data out of public AI tools.
- Decide whether the next best step is self-practice, one-on-one coaching, a team briefing, or a structured workshop.

Guiding questions

- What can an LLM help me draft, reframe, summarize, or plan?
- What information should stay out of public AI tools?
- How do I tell whether an AI answer is useful enough to keep working with?
- Which everyday task would be a safe first place to practice?

Readiness checks

- Can explain what an LLM is without technical jargon.
- Can write a prompt with context, task, audience, and constraints.
- Can name at least three kinds of data that should not be pasted into public tools.
- Can identify when an output needs outside verification or human approval.

Curriculum modules

LLMs in plain language

A beginner-friendly explanation of how large language models respond to prompts, what they are useful for, and where they become unreliable.

The first prompting loop

A reusable pattern for giving context, asking for a draft, checking the result, and requesting a revision without overtrusting the output.

Safe public-tool practice

A practical boundary-setting module that keeps sensitive or proprietary details out of public tools while still making room to learn.

Review before use

A lightweight review habit for checking accuracy, tone, assumptions, missing context, and decisions that need a human owner.

Choosing a next step

A closing module that helps learners identify one useful, low-risk task to practice and one question to bring to a trainer or manager.

Session flow

Orient

Start with a plain-language explanation of LLMs, prompts, outputs, and why confident-sounding answers still need review.

Try

Use a neutral sample task to practice context, constraints, tone, examples, and follow-up questions.

Review

Check the result for accuracy, usefulness, assumptions, missing context, and whether a human should verify it.

Choose

Name one safe workflow to practice and one boundary the learner will keep in mind.

Practice labs

Rewrite a neutral paragraph

Practice asking for tone and audience changes using public, non-sensitive text.

Summarize public information

Turn a public article or generic meeting note into a summary, then identify what still needs verification.

Plan a simple task

Ask AI for a checklist, timeline, or agenda, then revise it to fit real constraints.

Critique the output

Use a review prompt to find weak assumptions, vague wording, and missing context.

Tangible cases

Public article summary

Situation: A learner wants to understand a public article, newsletter, or announcement without losing the main ideas.

Learner task: Use AI to summarize the text, then mark which claims still need source checking before sharing.

Starter prompt: Summarize the following public text for a beginner. Give me five key points, three terms to define, and a short list of claims I should verify before I rely on the summary: [paste public text].

Fictional email draft

Situation: A professional wants to practice drafting an email without using client, patient, employee, or proprietary details.

Learner task: Draft from a fictional scenario, revise tone, and decide what a human should edit before sending.

Starter prompt: Using this fictional scenario, draft a warm professional email. Keep it concise, avoid promises, and include a checklist of details a human should confirm before sending: [fictional scenario].

First practice plan

Situation: A curious learner needs a low-pressure way to keep practicing after the starter session.

Learner task: Choose one safe weekly task, define boundaries, and keep a short log of useful and unreliable outputs.

Starter prompt: Help me design a one-week AI practice plan using only public or fictional information. Include one task per day, a safety reminder, and a simple log format for what worked, what failed, and what I checked.

Prompt library

Plain-language explanation

Explain [AI concept] for a smart beginner. Use an everyday analogy, name two useful applications, and name two reasons I should not overtrust the output.

Starter prompt builder

Turn my rough request into a stronger prompt. Ask me for missing context first. Then produce a prompt with task, audience, context, constraints, and review criteria.

Output review

Review this AI-generated draft. Identify unsupported claims, vague wording, missing context, possible bias, tone problems, and anything that needs human verification: [draft].

Safe practice chooser

Given these possible practice tasks, sort them into safe public practice, use caution, and do not use in a public tool. Explain the reason for each category: [task list].

Materials to develop

- Free AI fluency starter guide
- Basic prompting loop worksheet
- Public-tool safety checklist
- Output review checklist
- First safe workflow practice log

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- Beginner LLM session loop: ask, inspect, revise, verify
- Safe practice boundary map for public AI tools
- First workflow chooser: draft, summarize, plan, or critique

Follow-up work

- Keep a one-week practice log using only non-sensitive examples.
- Collect three questions that would benefit from coaching or a team discussion.
- Choose one recurring task that may become a future workflow workshop candidate.

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.