

Free AI Fluency Starter

A public starter path for people who want a clear, non-intimidating first step into LLMs, prompting, safe practice, and practical review habits.

Format	Self-guided webpage and downloadable PDF companion, designed as a gentle entry point before a team briefing or coached workshop.
Audience	Individuals, leaders, and teams who are curious about AI but want plain-language basics before committing to training.

Learner outcomes

- Understand what LLMs do, why they can be useful, and why their answers still need human review.
- Use a simple prompt structure for common tasks such as drafting, summarizing, planning, and asking better questions.
- Practice with non-sensitive examples while keeping patient, customer, employee, legal, financial, and proprietary data out of public AI tools.

Curriculum modules

LLMs in plain language

A beginner-friendly explanation of how large language models respond to prompts, what they are useful for, and where they become unreliable.

The first prompting loop

A reusable pattern for giving context, asking for a draft, checking the result, and requesting a revision without overtrusting the output.

Safe public-tool practice

A practical boundary-setting module that keeps sensitive or proprietary details out of public tools while still making room to learn.

Review before use

A lightweight review habit for checking accuracy, tone, assumptions, missing context, and decisions that need a human owner.

Materials to develop

- Free AI fluency starter guide
- Basic prompting loop worksheet
- Public-tool safety checklist
- Output review checklist

Diagram candidates

These are intentionally left as candidate visuals until diagram parameters are chosen.

- Beginner LLM session loop: ask, inspect, revise, verify
- Safe practice boundary map for public AI tools

Session design reminders

- Keep examples non-sensitive and realistic.
- Emphasize education, judgment, review, and careful experimentation.
- Do not promise legal, medical, compliance, cybersecurity, or automation outcomes.
- Name when a stronger tool, private environment, or policy decision is needed before proceeding.

API and tool boundary

This curriculum companion does not require an API call. Any secure password protection, private client portal, or live AI workflow exercise that transmits data to an external model should be reviewed before implementation.